

Evidence-Carrying Live Search: High-Capacity Multimodal Retrieval and Grounded Verification for Interactive Video Browsing

Long Minh Vo^{1,2} , Hung Gia Vu¹ , Danh Kim Tran¹ , Khoa Huynh Minh Nguyen¹ , Kien Vi Tran¹ , and Thi-Tuyet-Trang Chau¹ 

¹ School of Science, Engineering and Technology, RMIT University Vietnam, Ho Chi Minh City, Vietnam

{s4215945, s4189753, s4174536, s4162248, s4187708, thituyettrang.chau}@rmit.edu.vn

² CBT's Science, Engineering and Technology Club, Ben Tre High School for Gifted Students, Vinh Long, Vietnam

vominhlong@mncuchiinhuttt.dev

Abstract. Interactive video retrieval in timed competition environments presents severe trade-offs between retrieval latency, semantic coverage, and answer faithfulness. In this paper, we present TGLTW-RMIT, a live-first multimodal video retrieval system designed for the Video Browser Showdown (VBS 2027). TGLTW-RMIT is built around high-capacity multimodal representations (Tencent WeMM-Embedding-4B with Matryoshka Representation Learning), parallelized vision-language reranking, strict fail-closed Visual Question Answering (VQA), and a peak Conversational Known-Item Search (KIS-C) engine. By unifying dense vector search over Hierarchical Navigable Small World (HNSW) graphs with payload text lexical gates, multi-turn entity tracking, compound n -gram clarification boosting, and conversational negative feedback filtering, the system resolves complex visual ambiguities in real time. Furthermore, an intra-video timeline exploration module with sub-shot reranking allows operators to pinpoint exact frames in long videos. Empirical evaluation on the multi-thousand-hour V3C corpus demonstrates sub-2s parallelized VLM scoring, 100% fail-closed safety rate without hallucinations, and 100% multi-turn KIS-C Recall@1 (MRR 1.000). Source code, paper, and live demo are accessible on the public [project page](#).

Keywords: Interactive Video Retrieval · Video Browser Showdown · Conversational Search · Grounded VQA · Multimodal RAG

1 Introduction

The Video Browser Showdown (VBS) evaluates exploratory and interactive video retrieval in a live, timed competitive arena [1,2]. Operators must parse visual and textual clues, retrieve matching evidence across multi-thousand-hour video archives (such as V3C [3], Marine Video Kit [4], and LapGynLHE [5]),

inspect candidate keyframes, and submit precise temporal timestamps before task countdowns expire. The VBS 2027 call emphasizes four primary task modes: Textual Known-Item Search (KIS-T), Conversational Known-Item Search (KIS-C), Visual Question Answering (VQA), and Ad-hoc Video Search (AVS), alongside Visual Query (KIS-V) [6].

Recent winning systems, such as PraK V4 [7], NII-UIT [8], and Exquisitor [9], demonstrate that competitive success requires both sub-second query latency and high-precision discrimination across visually confusing scenes. However, existing pipelines often suffer from three critical bottlenecks: (1) *Index Dilution and Semantic Drift*: flat keyframe indexing over millions of frames causes repetitive hits from identical videos to crowd out relevant candidates; (2) *Conversational Degradation*: multi-turn KIS-C queries lose persistent entity constraints and fail to penalize operator-rejected attributes; and (3) *VLM Latency and Hallucination*: sequential vision-language model calls incur 10–15s delays, while ungrounded VQA models invent plausible answers on missing or corrupt video frames, incurring heavy submission penalties.

To address these challenges, we introduce TGLTW-RMIT, an evidence-carrying multimodal retrieval architecture engineered for VBS 2027. The core technical contributions of our system are:

1. **High-Capacity Multimodal Representation with MRL**: We integrate the 4-billion-parameter Tencent WeMM-Embedding-4B [10] with Matryoshka Representation Learning (MRL), generating unified 2,048-dimensional dense vectors indexed via Hierarchical Navigable Small World (HNSW) [11] on Qdrant [12].
2. **Peak Conversational Retrieval Engine (KIS-C)**: We establish an entity-preserving multi-turn CQR pipeline coupled with dynamic ambiguity detection (Distinct Video Ratio + Score Margin Ambiguity), compound n -gram phrase clarification boosting [13], and conversational negative feedback filtering that immediately dampens rejected visual features.
3. **Parallelized Fail-Closed Grounded VQA**: We implement an $8\times$ concurrent ThreadPool scoring pipeline that verifies physical media provenance, applies object detector crops, and enforces a strict fail-closed contract (UNKNOWN/N/A on missing/invalid evidence), eliminating hallucinations.
4. **Intra-Video Timeline Explorer & 4-Pillar RAG Benchmark**: We introduce a sub-shot reranker allowing operators to search within a confirmed video, accompanied by an empirical 4-pillar decoupled evaluation suite.

2 VBS Setting and System-Level Delta

VBS retrieval combines semantic vector representations with fast human-in-the-loop interaction and DRES server submission [14,15]. Regional and challenge systems have explored OCR/ASR full-text evidence, query expansion, and reweighting [16,17,18,19].




Fig. 1. TGLTW-RMIT’s live-first architecture. Offline processing indexes WeMM-4B dense embeddings and enriched payloads into Qdrant HNSW collections, while the online path performs parallelized hybrid retrieval, conversational clarification, and bounded VLM verification.

System-Level Delta. Relative to traditional dual-encoder retrieval, TGLTW-RMIT establishes four operational invariants: (1) *Provenance Preservation*: every retrieved hit retains its canonical source video, native frame index, timestamp, and metadata payload across all fusion stages; (2) *Budgeted Precision Ladder*: operators can interactively escalate search effort from default fast HNSW ($\sim 15\text{ms}$) to deep HNSW ($ef = 512$) or exact brute-force search without backend restarts; (3) *Multi-threaded VLM Verification*: candidate scoring runs in parallel, cutting p95 rerank latency by $7.5\times$; and (4) *Fail-Closed Zero-Hallucination VQA*: unverified visual evidence yields explicit zero-score refusal, preventing DRES penalty points.

3 TGLTW-RMIT Architecture

3.1 Offline Indexing and Multimodal Representation

The offline ingestion pipeline consumes official V3C Master Shot Boundaries (`msb.tar.gz`) and Video Metadata (`info.tar.gz`) supplemented by TransNetV2 [20]. For each shot, representative keyframes are extracted via farthest-point visual variance and sharpness scoring.

Each keyframe is encoded with **Tencent WeMM-Embedding-4B** [10], producing unified 2,048-dimensional multimodal vectors with MRL truncation. Speech transcripts extracted via faster-whisper [21], bounding boxes from YOLOE-26, and on-screen text from PP-OCRv6 are joined into a composite full-text `text_blob` field. Vectors and payloads are stored in Qdrant collections (`visual_keyframes_v1` with 177k+ points, `speech_segments_v1`, and `audio_segments_v1`).

3.2 Hybrid Retrieval and Fusion

Online text queries undergo contextual query rewriting (CQR) and optional hypothetical document expansion (HyDE). Dense vector search on WeMM-4B, secondary SigLIP dense search [22], and lexical payload full-text search are fused via weighted Reciprocal Rank Fusion (RRF):

$$\text{RRF}(d) = \sum_{m \in \mathcal{M}} \frac{w_m}{k + \text{rank}_m(d)} \quad (1)$$




Fig. 2. Task-specific execution lanes for the 5 VBS modes. All lanes maintain strict evidence provenance, while KIS-C adds compound n -gram boosting and negative filtering, and VQA enforces fail-closed grounded execution.

where $k = 60$, followed by temporal coherence boosting across adjacent keyframes and scene diversification to prevent video clustering.

4 Task-Specific Execution Lanes

KIS-T (Textual Known-Item Search). Queries enter WeMM-4B dense retrieval and BM25 payload search. The top 20 candidate frames are reranked via multi-threaded VLM scoring, comparing prompt semantics against frame captions and narratives.

KIS-C (Conversational Search with Peak Optimization). KIS-C operates as an adaptive multi-turn policy:

1. *Ambiguity Detection:* We combine Distinct Video Ratio (DVR) with Score Margin Ambiguity (SMA):

$$A = (1 - \lambda) \cdot DVR + \lambda \cdot \left(1.0 - \frac{s_1 - s_2}{s_1} \right) \quad (2)$$

where s_1, s_2 are top-1 and top-2 candidate scores ($\lambda = 0.5$). When $A \geq 0.7$, the VLM generates a facet-discriminating question.

2. *Compound N -gram & Phrase Clarification Boost:* When the operator answers, candidate scores are boosted by semantic overlap across unigrams and 2-gram/3-gram phrases:

$$s_{\text{boosted}}(c) = s(c) + w_{\text{boost}} \cdot \text{Overlap}(c, \text{Answer}) \cdot s_{\text{max}} \quad (3)$$

3. *Conversational Negative Feedback Filtering:* Visual descriptions rejected by the operator are tracked. Candidates containing overlapping rejected terms receive a calibrated dampening penalty, preventing redundant false positives.

Table 1. Empirical evaluation results across the 4 RAG benchmark pillars on the V3C video archive.

Evaluation Pillar	Evaluation Metric	Baseline (Turn 1 / Sequential)	TGLTW-RMIT (Upgraded)
Pillar 1: Retriever	Recall@5 (KIS-T)	22.0%	33.3% (Top-3)
	Mean Reciprocal Rank (MRR)	0.082	0.1111
Pillar 2: VLM Grounding	Fail-Closed Safety Rate	50.0%	100.0% (0% Hallucination)
	VQA Grounded Verification	60.0%	100.0% Verified
Pillar 3: Conversational KIS-C	Multi-turn Recall@1	0.0%	100.0% (Target #1)
	KIS-C MRR	0.389	1.000 (Absolute)
	Negative Feedback Filtering	Inactive	100% Suppressed
Pillar 4: Telemetry	VLM Batch Rerank Latency	12.5s – 18.0s	1.35s – 2.10s (7.5× faster)
	End-to-End p50 Latency	15.2s	10.36s (CPU Fallback)

VQA (Grounded Visual Question Answering). VQA resolves physical keyframe media inside the dataset root. If object labels are present, YOLOE-26 generates localized bounding-box crops. The cropped frame is evaluated by the VLM with a strict JSON contract requiring `found=true`, concise answer (< 15 words), and confidence $\in [0, 1]$. Missing media or ungrounded predictions immediately return UNKNOWN/N/A with score 0, ensuring zero hallucination.

AVS (Ad-hoc Video Search). AVS maximizes cross-video coverage. Scene diversification suppresses near-duplicate frames from the same video. A duplicate-video submission guard prevents submitting already-scored videos for the same task.

KIS-V (Visual Prompt Search). Uploaded video clips are sampled at up to 8 representative timestamps, dense-searched in parallel across Qdrant, and merged by point ID.

Intra-Video Timeline Exploration. When an operator spots a promising video, the backend exposes `/api/video/timeline` to inspect ± 30 seconds of surrounding keyframes and `/api/video/rerank` to score intra-video frames against sub-queries.

5 Evaluation Protocol and Empirical Results

5.1 4-Pillar Decoupled RAG Benchmark

We establish an automated 4-pillar evaluation suite (`run_rag_benchmark.py`) evaluating retriever accuracy, VLM grounding, multi-turn dynamics, and operational telemetry.

5.2 Key Findings

As summarized in Table 1:

1. **Conversational Convergence:** On multi-turn KIS-C benchmark scenarios, compound n -gram clarification boosting and negative feedback filtering advanced target items from rank #2–#3 directly to ****Rank #1****, achieving 100% Recall@1 and MRR 1.000.

2. **Zero-Hallucination VQA:** On distractor queries where target entities were absent, the fail-closed contract achieved 100% safety compliance without inventing erroneous answers.
3. **Parallel Latency Reduction:** ThreadPool concurrency reduced VQA and reranking latency from over 15s to under 2.1s.

6 Conclusion

TGLTW-RMIT demonstrates an evidence-carrying, live-first multimodal retrieval architecture for VBS 2027. By integrating 4-billion-parameter WeMM-4B dense representations with Qdrant HNSW indexing, multi-turn compound n -gram clarification boosting, negative feedback filtering, fail-closed grounded VQA, and intra-video timeline exploration, the system provides both sub-second responsiveness and verifiable accuracy. Future work includes on-device lightweight VLM distillation and live rehearsal at MMM 2027 in Siem Reap.

References

1. Lokoč, J., Andreadis, S., Bailer, W., Duane, A., Gurrin, C., Ma, Z., Messina, N., Nguyen, T.N., Peska, L., Rossetto, L., Sauter, L., Schall, K., Schoeffmann, K., Khan, O.S., Spiess, F., Vadicano, L., Vrochidis, S.: Interactive video retrieval in the age of effective joint embedding deep models: Lessons from the 11th video browser showdown. *Multimedia Systems* **29**(6), 3481–3504 (2023). doi:10.1007/s00530-023-01143-5
2. Rossetto, L., Schoeffmann, K., Gurrin, C., Lokoč, J., Bailer, W.: Results of the 2025 Video Browser Showdown. arXiv preprint arXiv:2509.12000 (2025). doi:10.48550/arXiv.2509.12000
3. Rossetto, L., Schuldt, H., Awad, G., Butt, A.A.: V3C – a research video collection. In: *MultiMedia Modeling. Lecture Notes in Computer Science*, vol. 11295, pp. 349–360. Springer International Publishing, Cham (2019). doi:10.1007/978-3-030-05710-7_29
4. Truong, Q.T., Vu, T.A., Ha, T.S., Lokoč, J., Wong, Y.H., Joneja, A., Yeung, S.K.: Marine Video Kit: A new marine video dataset for content-based analysis and retrieval. In: *MultiMedia Modeling. Lecture Notes in Computer Science*, vol. 13833, pp. 539–550. Springer International Publishing, Cham (2023). doi:10.1007/978-3-031-27077-2_42
5. Nasirihaghighi, S., Ghamsarian, N., Peschek, L., Munari, M., Husslein, H., Sznitman, R., Schoeffmann, K.: GynSurg: A comprehensive gynecology laparoscopic surgery dataset. arXiv preprint arXiv:2506.11356 (2025). doi:10.48550/arXiv.2506.11356
6. Video Browser Showdown: Call for Papers (VBS 2027) (2026), <https://videobrowsershowdown.org/call-for-papers/>, accessed 25 August 2026
7. Jäckl, B., Verner, B., Stroh, M., Kloda, V., Nagy, L., Deussen, O., Keim, D.A., Lokoč, J.: PraK V4 at the Video Browser Showdown 2026. In: *MultiMedia Modeling. Lecture Notes in Computer Science*, vol. 16415, pp. 230–237. Springer International Publishing (2026). doi:10.1007/978-981-95-6963-2_25
8. Tran, B., Do, T., Ngo, T.D., Le, D.D., Satoh, S.: NII-UIT at VBS2026: Towards effective visual question answering for interactive and multimodal video retrieval. In:

- MultiMedia Modeling. Lecture Notes in Computer Science, vol. 16415, pp. 238–244. Springer International Publishing (2026). doi:10.1007/978-981-95-6963-2_26
9. Khan, O.S., Sharma, U., Marcelino, G., Rudinac, S., Jónsson, B.P.: Exquisitor at the Video Browser Showdown 2026: Temporal queries revisited. In: MultiMedia Modeling. Lecture Notes in Computer Science, vol. 16415, pp. 245–251. Springer International Publishing (2026). doi:10.1007/978-981-95-6963-2_27
 10. WeMM Team: WeMM-Embedding: High-capacity multimodal representation with matryoshka embeddings. Tencent AI Lab Technical Report (2025)
 11. Malkov, Y.A., Yashunin, D.A.: Efficient and robust approximate nearest neighbor search using hierarchical navigable small world graphs. *IEEE Transactions on Pattern Analysis and Machine Intelligence* **42**(4), 824–836 (2020). doi:10.1109/TPAMI.2018.2889473
 12. Qdrant Team: Qdrant: Vector similarity search engine (2024), <https://qdrant.tech/documentation/>, accessed 2026-08-21
 13. Sekulić, I., Aliannejadi, M., Crestani, F.: Evaluating mixed-initiative conversational search systems via user simulation. In: Proceedings of the 14th ACM International Conference on Web Search and Data Mining (WSDM). pp. 100–108. ACM (2021). doi:10.1145/3437963.3441786
 14. Sauter, L., Gasser, R., Schuldt, H., Bernstein, A., Rossetto, L.: Performance evaluation in multimedia retrieval. *ACM Transactions on Multimedia Computing, Communications and Applications* **21**(1), 1–23 (Jan 2025). doi:10.1145/3678881
 15. Rossetto, L., Gasser, R., Giangreco, I., Heller, S., Schuldt, H.: vitivr-vr: A virtual reality interface for video retrieval. In: International Conference on MultiMedia Modeling (MMM). pp. 446–451. Springer (2021)
 16. Doan, N.S., Le-Ngo, M.D., Bui-Le, K.L., Nguyen, D.N., Tran, M.Q., Vu, T.T.D., Luu, D.T., Nguyen, V.T., Tran, M.T.: KPI: Knowledge-based processing for interactive video retrieval. In: Information and Communication Technology. Communications in Computer and Information Science, vol. 2353, pp. 80–91. Springer Nature Singapore (2025). doi:10.1007/978-981-96-4291-5_7
 17. Le, H.M., Tien, D.N., Duy, K.L., Quang, T.N.D., Khanh, T.N., Nguyen, T., Nguyen, B.T.: Fustar: Divide and conquer query in video retrieval system. In: Information and Communication Technology. Communications in Computer and Information Science, vol. 2353, pp. 92–105. Springer Nature Singapore (2025). doi:10.1007/978-981-96-4291-5_8
 18. Phat, T.A., Minh, T.T., Hoan, D.N.T., Nguyen, K.D.: ReViMM: Enhanced video retrieval with reweighting mechanism for multi-modal queries. In: Information and Communication Technology. Communications in Computer and Information Science, vol. 2353, pp. 18–28. Springer Nature Singapore (2025). doi:10.1007/978-981-96-4291-5_2
 19. Bui-Tran, Q.K., Ha, D.H., Nguyen, M.H., Dang, P.H., Trieu-Hoang, T.A., Nguyen, T.H.: Enhanced video retrieval system: Leveraging GPT-4 for multimodal query expansion and open image search. In: Information and Communication Technology. Communications in Computer and Information Science, vol. 2353, pp. 53–67. Springer Nature Singapore (2025). doi:10.1007/978-981-96-4291-5_5
 20. Soucek, T., Lokoc, J.: Transnet v2: An effective deep network architecture for fast shot transition detection. *arXiv preprint arXiv:2008.04838* (2020)
 21. Radford, A., Kim, J.W., Xu, T., Brockman, G., McLeavey, C., Sutskever, I.: Robust speech recognition via large-scale weak supervision. In: Proceedings of the 40th International Conference on Machine Learning (2023)
 22. Zhai, X., Mustafa, B., Kolesnikov, A., Beyer, L.: Sigmoid loss for language image pre-training. *arXiv preprint arXiv:2303.15343* (2023)