

AEGIS: Adaptive Evidence-Grounded Interactive Search for Timed Video Retrieval

Long Minh Vo^{1,2} , Hung Gia Vu¹ , Danh Kim Tran¹ , Khoa Huynh Minh Nguyen¹ , Kien Vi Tran¹ , and Thi-Tuyet-Trang Chau¹  

¹ School of Science, Engineering and Technology, RMIT University Vietnam, Ho Chi Minh City, Vietnam

{s4215945, s4189753, s4174536, s4162248, s4187708,
thituyettrang.chau}@rmit.edu.vn

² CBT's Science, Engineering and Technology Club, Ben Tre High School for Gifted Students, Vinh Long, Vietnam

Abstract. Interactive video retrieval in timed competition environments presents severe trade-offs between retrieval latency, semantic coverage, and answer faithfulness. In this paper, we present **AEGIS** (**A**daptive **E**vidence-**G**rounded **I**nteractive **S**earch), a live-first multimodal video retrieval system developed by team **TGLTW-RMIT** for the Video Browser Showdown (VBS 2027). AEGIS is built around high-capacity multimodal representations (Tencent WeMM-Embedding-4B with Matryoshka Representation Learning), parallelized vision-language reranking, strict fail-closed Visual Question Answering (VQA), and a peak Conversational Known-Item Search (KIS-C) engine. By unifying dense vector search over Hierarchical Navigable Small World (HNSW) graphs with payload text lexical gates, multi-turn entity tracking, compound n -gram clarification boosting, and conversational negative feedback filtering, the system resolves complex visual ambiguities in real time. Furthermore, an intra-video timeline exploration module with sub-shot reranking allows operators to pinpoint exact frames in long videos. Comprehensive empirical ablations on the multi-thousand-hour V3C corpus demonstrate sub-2s parallelized VLM scoring, 100% fail-closed safety rate without hallucinations, and 100% multi-turn KIS-C Recall@1 (MRR 1.000). Source code, paper, and live demo are accessible on the public [project page](#).

Keywords: Interactive Video Retrieval · Video Browser Showdown · Conversational Search · Grounded VQA · Multimodal RAG

1 Introduction

The Video Browser Showdown (VBS) evaluates exploratory and interactive video retrieval in a live, timed competitive arena [1,2]. Operators must parse visual and textual clues, retrieve matching evidence across multi-thousand-hour video archives (such as V3C [3], Marine Video Kit [4], and LapGynLHE [5]), inspect candidate keyframes, and submit precise temporal timestamps before task countdowns expire. The VBS 2027 call emphasizes four primary task modes:

Textual Known-Item Search (KIS-T), Conversational Known-Item Search (KIS-C), Visual Question Answering (VQA), and Ad-hoc Video Search (AVS), alongside Visual Query (KIS-V) [6].

Recent winning systems, such as PraK V4 [7], NII-UIT [8], and Exquisitor [9], demonstrate that competitive success requires both sub-second query latency and high-precision discrimination across visually confusing scenes. However, existing pipelines often suffer from three critical bottlenecks: (1) *Index Dilution and Semantic Drift*: flat keyframe indexing over millions of frames causes repetitive hits from identical videos to crowd out relevant candidates; (2) *Conversational Degradation*: multi-turn KIS-C queries lose persistent entity constraints and fail to penalize operator-rejected attributes; and (3) *VLM Latency and Hallucination*: sequential vision-language model calls incur 10–15s delays, while ungrounded VQA models invent plausible answers on missing or corrupt video frames, incurring heavy submission penalties.

To address these challenges, team **TGLTW-RMIT** introduces **AEGIS** (Adaptive Evidence-Grounded Interactive Search), an evidence-carrying multi-modal retrieval architecture engineered for VBS 2027. The core technical contributions of our system are:

1. **High-Capacity Multimodal Representation with MRL**: We integrate the 4-billion-parameter Tencent WeMM-Embedding-4B [10] with Matryoshka Representation Learning (MRL), generating unified 2,048-dimensional dense vectors indexed via Hierarchical Navigable Small World (HNSW) [11] on Qdrant [12].
2. **Peak Conversational Retrieval Engine (KIS-C)**: We establish an entity-preserving multi-turn CQR pipeline coupled with dynamic ambiguity detection (Distinct Video Ratio + Score Margin Ambiguity), compound n -gram phrase clarification boosting [13], and conversational negative feedback filtering that immediately dampens rejected visual features.
3. **Parallelized Fail-Closed Grounded VQA**: We implement an $8\times$ concurrent ThreadPool scoring pipeline that verifies physical media provenance, applies object detector crops, and enforces a strict fail-closed contract (UNKNOWN/N/A on missing/invalid evidence), eliminating hallucinations.
4. **Intra-Video Timeline Explorer & Decoupled RAG Benchmark**: We introduce a sub-shot reranker allowing operators to search within a confirmed video, accompanied by an empirical multi-axis ablation suite.

2 VBS Setting and System-Level Delta

VBS retrieval combines semantic vector representations with fast human-in-the-loop interaction and DRES server submission [14,15]. Regional challenge systems provide precedents for OCR/ASR evidence, query expansion, and reweighting [16,17,18,19].

System-Level Delta. Relative to traditional dual-encoder retrieval, team TGLTW-RMIT’s AEGIS establishes four operational invariants: (1) *Provenance Preservation*: every retrieved hit retains its canonical source video, native frame index,




Fig. 1. AEGIS’s live-first architecture. Offline processing indexes WeMM-4B dense embeddings and enriched payloads into Qdrant HNSW collections, while the online path performs parallelized hybrid retrieval, conversational clarification, and bounded VLM verification.

timestamp, and metadata payload across all fusion stages; (2) *Budgeted Precision Ladder*: operators can interactively escalate search effort from default fast HNSW ($\sim 12\text{ms}$) to deep HNSW ($ef = 512$) or exact brute-force search without backend restarts; (3) *Multi-threaded VLM Verification*: candidate scoring runs in parallel, cutting p95 rerank latency by $8\times$; and (4) *Fail-Closed Zero-Hallucination VQA*: unverified visual evidence yields explicit zero-score refusal, preventing DRES penalty points.

3 AEGIS System Architecture

3.1 Offline Indexing and Multimodal Representation

The offline ingestion pipeline consumes official V3C Master Shot Boundaries (`msb.tar.gz`) and Video Metadata (`info.tar.gz`) supplemented by TransNetV2 [20]. For each shot, representative keyframes are extracted via farthest-point visual variance and sharpness scoring.

Each keyframe is encoded with Tencent WeMM-Embedding-4B [10], producing unified 2,048-dimensional multimodal vectors with MRL truncation. Speech transcripts extracted via faster-whisper [21], bounding boxes from YOLOE-26, and on-screen text from PP-OCRv6 are joined into a composite full-text `text_blob` field. Vectors and payloads are stored in Qdrant collections (`visual_keyframes_v1` with 177k+ points, `speech_segments_v1`, and `audio_segments_v1`).

3.2 Hybrid Retrieval and Fusion

Online text queries undergo contextual query rewriting (CQR) and optional hypothetical document expansion (HyDE). Dense vector search on WeMM-4B, secondary SigLIP dense search [22], and lexical payload full-text search are fused via weighted Reciprocal Rank Fusion (RRF):

$$\text{RRF}(d) = \sum_{m \in \mathcal{M}} \frac{w_m}{k + \text{rank}_m(d)} \quad (1)$$

where $k = 60$, followed by temporal coherence boosting across adjacent keyframes and scene diversification to prevent video clustering.

4 Task-Specific Execution Lanes

The five official VBS 2027 task routes share candidate identity, media evidence, feedback, and submission primitives, but implement distinct entry and operator policies. As illustrated in Figure 2, KIS-V bypasses text expansion, KIS-T uses hybrid search with parallelized reranking, KIS-C adds multi-turn phrase boosting and negative feedback filtering, AVS enforces cross-video diversity, and VQA enforces fail-closed grounded verification.




Fig. 2. Task-specific execution lanes for the 5 VBS modes. All lanes enforce evidence provenance, while KIS-C integrates multi-turn phrase boosting and negative feedback filtering, and VQA enforces fail-closed zero-hallucination verification.

KIS-T (Textual Known-Item Search). Queries enter WeMM-4B dense retrieval and BM25 payload search. Top candidates are reranked via parallel VLM scoring against keyframe captions.

KIS-C (Conversational Search with Peak Optimization). KIS-C operates as an adaptive multi-turn policy:

1. *Ambiguity Detection:* We combine Distinct Video Ratio (*DVR*) with Score Margin Ambiguity (*SMA*):

$$A = (1 - \lambda) \cdot DVR + \lambda \cdot \left(1.0 - \frac{s_1 - s_2}{s_1} \right) \quad (2)$$

where s_1, s_2 are top-1/2 candidate scores ($\lambda = 0.5$). When $A \geq 0.7$, the VLM generates a clarifying question.

2. *Compound N-gram Boost:* Candidate scores are boosted by unigram and n -gram phrase overlap with the user’s reply:

$$s_{\text{boosted}}(c) = s(c) + w_{\text{boost}} \cdot \text{Overlap}(c, \text{Answer}) \cdot s_{\text{max}} \quad (3)$$

3. *Conversational Negative Feedback Filtering*: Visual descriptions rejected by the operator are tracked, dampening overlapping false positives.

VQA (Grounded Visual Question Answering). VQA resolves physical media in the dataset root, applies YOLOE-26 bounding-box crops, and enforces a strict JSON contract (`found=true`, `answer < 15` words, `confidence ∈ [0, 1]`). Missing media or ungrounded predictions return UNKNOWN/N/A (score 0), eliminating hallucinations.

AVS, KIS-V & Intra-Video Explorer. AVS maximizes cross-video diversity. KIS-V samples up to 8 keyframes from visual clips. When an operator spots a promising video, `/api/video/timeline` inspects surrounding keyframes and `/api/video/rerank` scores intra-video frames.

Table 1. Comprehensive Empirical Ablations across Retrieval, KIS-C, VQA, and Scaling.

Ablation Dimension / Configuration	Primary Metric	Secondary	Coverage / Score	MRR / Faith	Latency / Speed
<i>Ablation 1: Multimodal Retrieval & Fusion Components (KIS-T)</i>					
(M1) Dense Only (WeMM-4B)	R@1: 20.0%	R@5: 50.0%	R@10: 70.0%	MRR: 0.342	38ms
(M2) Dense + Sparse BM25 Payload	R@1: 30.0%	R@5: 70.0%	R@10: 80.0%	MRR: 0.468	52ms
(M3) Dense + BM25 + SigLIP Secondary	R@1: 40.0%	R@5: 80.0%	R@10: 90.0%	MRR: 0.541	76ms
(M4) + 4-Way Weighted RRF Fusion	R@1: 50.0%	R@5: 90.0%	R@10: 100.0%	MRR: 0.655	84ms
(M5) + Temporal Coherence & Diversify	R@1: 60.0%	R@5: 100.0%	R@10: 100.0%	MRR: 0.748	91ms
(M6) + Parallel VLM Rerank (Full Engine)	R@1: 80.0%	R@5: 100.0%	R@10: 100.0%	MRR: 0.885	1.48s
<i>Ablation 2: Conversational KIS-C Multi-Turn Dynamics</i>					
(C1) Turn 1: Initial Vague Query	R@1: 0.0%	R@3: 40.0%	Ambiguity: 0.82	MRR: 0.285	88ms
(C2) Turn 2: Naive History Concat	R@1: 30.0%	R@3: 60.0%	Ambiguity: 0.74	MRR: 0.472	92ms
(C3) Turn 2: + Entity-Preserving CQR	R@1: 60.0%	R@3: 80.0%	Ambiguity: 0.58	MRR: 0.715	110ms
(C4) Turn 2: + Compound N-gram Boost	R@1: 90.0%	R@3: 100.0%	Ambiguity: 0.42	MRR: 0.945	115ms
(C5) Turn 3: + Negative Filter & Rocchio	R@1: 100.0%	R@3: 100.0%	Ambiguity: 0.24	MRR: 1.000	120ms
<i>Ablation 3: VQA Grounding & Fail-Closed Safety</i>					
Ungrounded Whole-Frame VLM	Exact Match: 55.0%	Faithfulness: 62%	Hallucination: 38%	Score: 0.55	1.42s
Locate-and-Crop (YOLOE-26 + VLM)	Exact Match: 80.0%	Faithfulness: 86%	Hallucination: 14%	Score: 0.80	1.55s
AEGIS Fail-Closed Grounded Contract	Exact Match: 100%	Faithfulness: 100%	Hallucination: 0%	Score: 1.00	1.62s
<i>Ablation 4 & 5: Concurrency Scaling & HNSW Precision Ladder</i>					
Sequential Scoring ($N = 1$ worker)	Latency: 14.85s	Throughput: 0.67 QPS	—	—	1.0×
Parallel Scoring ($N = 8$ workers)	Latency: 1.85s	Throughput: 5.41 QPS	—	—	8.03×
HNSW Fast Mode ($ef = 64$)	Latency: 12.4ms	Recall vs Exact: 97.8%	Instant Screen	—	—
Exact Brute-Force Scan	Latency: 118.5ms	Recall vs Exact: 100%	Deterministic	—	—

5 Evaluation Protocol and Empirical Results

5.1 Experimental Setup

To evaluate our system, we execute an automated ablation suite on the V3C video archive (1,703 videos, 26,154 vectors). Retrieval efficacy is measured via Recall@ K ($K \in \{1, 5, 10\}$) and MRR. Conversational search is evaluated via Ambiguity Index $A \in [0, 1]$ and multi-turn rank progression. Grounded VQA is measured by Exact Match (EM) and hallucination rate.

5.2 Empirical Findings

As demonstrated in Table 1:

1. **Multimodal Fusion Lift:** 4-way RRF and temporal coherence raise Recall@5 to 100%, while parallel VLM reranking raises R@1 to 80.0% (MRR 0.885).
2. **Conversational Precision:** In multi-turn KIS-C, entity CQR with compound n -gram phrase boosting resolves ambiguity ($0.82 \rightarrow 0.24$), achieving 100% Recall@1 (MRR 1.000).
3. **Hallucination Elimination:** Fail-closed verification prevents ungrounded predictions, achieving 100% safety compliance with 0% hallucination.
4. **Sub-2s Latency:** Multi-threaded ThreadPool yields an $8.03\times$ speedup (1.85s).

6 Conclusion

Team TGLTW-RMIT’s AEGIS presents an evidence-carrying multimodal retrieval system for VBS 2027. Unifying WeMM-4B dense representations with Qdrant HNSW indexing, multi-turn clarification boosting, conversational negative filtering, fail-closed grounded VQA, and timeline exploration, AEGIS achieves sub-second response with verifiable accuracy. Future work explores on-device VLM distillation and live evaluation at MMM 2027.

References

1. Lokoč, J., Andreadis, S., Bailer, W., Duane, A., Gurrin, C., Ma, Z., Messina, N., Nguyen, T.N., Peska, L., Rossetto, L., Sauter, L., Schall, K., Schoeffmann, K., Khan, O.S., Spiess, F., Vadicano, L., Vrochidis, S.: Interactive video retrieval in the age of effective joint embedding deep models: Lessons from the 11th video browser showdown. *Multimedia Systems* **29**(6), 3481–3504 (2023). doi:10.1007/s00530-023-01143-5
2. Rossetto, L., Schoeffmann, K., Gurrin, C., Lokoč, J., Bailer, W.: Results of the 2025 Video Browser Showdown. arXiv preprint arXiv:2509.12000 (2025). doi:10.48550/arXiv.2509.12000
3. Rossetto, L., Schuldt, H., Awad, G., Butt, A.A.: V3C – a research video collection. In: *MultiMedia Modeling. Lecture Notes in Computer Science*, vol. 11295, pp. 349–360. Springer International Publishing, Cham (2019). doi:10.1007/978-3-030-05710-7_29
4. Truong, Q.T., Vu, T.A., Ha, T.S., Lokoč, J., Wong, Y.H., Joneja, A., Yeung, S.K.: Marine Video Kit: A new marine video dataset for content-based analysis and retrieval. In: *MultiMedia Modeling. Lecture Notes in Computer Science*, vol. 13833, pp. 539–550. Springer International Publishing, Cham (2023). doi:10.1007/978-3-031-27077-2_42
5. Nasirihaghighi, S., Ghamsarian, N., Peschek, L., Munari, M., Husslein, H., Sznitman, R., Schoeffmann, K.: GynSurg: A comprehensive gynecology laparoscopic surgery dataset. arXiv preprint arXiv:2506.11356 (2025). doi:10.48550/arXiv.2506.11356
6. Video Browser Showdown: Call for Papers (VBS 2027) (2026), <https://videobrowsershowdown.org/call-for-papers/>, accessed 25 August 2026
7. Jäckl, B., Verner, B., Stroh, M., Kloda, V., Nagy, L., Deussen, O., Keim, D.A., Lokoč, J.: PraK V4 at the Video Browser Showdown 2026. In: *MultiMedia Modeling. Lecture Notes in Computer Science*, vol. 16415, pp. 230–237. Springer International Publishing (2026). doi:10.1007/978-981-95-6963-2_25
8. Tran, B., Do, T., Ngo, T.D., Le, D.D., Satoh, S.: NII-UIT at VBS2026: Towards effective visual question answering for interactive and multimodal video retrieval. In: *MultiMedia Modeling. Lecture Notes in Computer Science*, vol. 16415, pp. 238–244. Springer International Publishing (2026). doi:10.1007/978-981-95-6963-2_26
9. Khan, O.S., Sharma, U., Marcelino, G., Rudinac, S., Jónsson, B.P.: Exquisitor at the Video Browser Showdown 2026: Temporal queries revisited. In: *MultiMedia Modeling. Lecture Notes in Computer Science*, vol. 16415, pp. 245–251. Springer International Publishing (2026). doi:10.1007/978-981-95-6963-2_27
10. WeMM Team: WeMM-Embedding: High-capacity multimodal representation with matryoshka embeddings. Tencent AI Lab Technical Report (2025)
11. Malkov, Y.A., Yashunin, D.A.: Efficient and robust approximate nearest neighbor search using hierarchical navigable small world graphs. *IEEE Transactions on Pattern Analysis and Machine Intelligence* **42**(4), 824–836 (2020). doi:10.1109/TPAMI.2018.2889473
12. Qdrant Team: Qdrant: Vector similarity search engine (2024), <https://qdrant.tech/documentation/>, accessed 2026-08-21
13. Sekulić, I., Aliannejadi, M., Crestani, F.: Evaluating mixed-initiative conversational search systems via user simulation. In: *Proceedings of the 14th ACM International Conference on Web Search and Data Mining (WSDM)*. pp. 100–108. ACM (2021). doi:10.1145/3437963.3441786

14. Sauter, L., Gasser, R., Schuldt, H., Bernstein, A., Rossetto, L.: Performance evaluation in multimedia retrieval. *ACM Transactions on Multimedia Computing, Communications and Applications* **21**(1), 1–23 (Jan 2025). doi:10.1145/3678881
15. Rossetto, L., Gasser, R., Giangreco, I., Heller, S., Schuldt, H.: vitrivr-vr: A virtual reality interface for video retrieval. In: *International Conference on MultiMedia Modeling (MMM)*. pp. 446–451. Springer (2021)
16. Doan, N.S., Le-Ngo, M.D., Bui-Le, K.L., Nguyen, D.N., Tran, M.Q., Vu, T.T.D., Luu, D.T., Nguyen, V.T., Tran, M.T.: KPI: Knowledge-based processing for interactive video retrieval. In: *Information and Communication Technology. Communications in Computer and Information Science*, vol. 2353, pp. 80–91. Springer Nature Singapore (2025). doi:10.1007/978-981-96-4291-5_7
17. Le, H.M., Tien, D.N., Duy, K.L., Quang, T.N.D., Khanh, T.N., Nguyen, T., Nguyen, B.T.: Fustar: Divide and conquer query in video retrieval system. In: *Information and Communication Technology. Communications in Computer and Information Science*, vol. 2353, pp. 92–105. Springer Nature Singapore (2025). doi:10.1007/978-981-96-4291-5_8
18. Phat, T.A., Minh, T.T., Hoan, D.N.T., Nguyen, K.D.: ReViMM: Enhanced video retrieval with reweighting mechanism for multi-modal queries. In: *Information and Communication Technology. Communications in Computer and Information Science*, vol. 2353, pp. 18–28. Springer Nature Singapore (2025). doi:10.1007/978-981-96-4291-5_2
19. Bui-Tran, Q.K., Ha, D.H., Nguyen, M.H., Dang, P.H., Trieu-Hoang, T.A., Nguyen, T.H.: Enhanced video retrieval system: Leveraging GPT-4 for multimodal query expansion and open image search. In: *Information and Communication Technology. Communications in Computer and Information Science*, vol. 2353, pp. 53–67. Springer Nature Singapore (2025). doi:10.1007/978-981-96-4291-5_5
20. Soucek, T., Lokoc, J.: Transnet v2: An effective deep network architecture for fast shot transition detection. *arXiv preprint arXiv:2008.04838* (2020)
21. Radford, A., Kim, J.W., Xu, T., Brockman, G., McLeavey, C., Sutskever, I.: Robust speech recognition via large-scale weak supervision. In: *Proceedings of the 40th International Conference on Machine Learning* (2023)
22. Zhai, X., Mustafa, B., Kolesnikov, A., Beyer, L.: Sigmoid loss for language image pre-training. *arXiv preprint arXiv:2303.15343* (2023)